

Interpretable Control for Unstable Systems via SINDy-RL

Patrick Smith · Andrew Falcone

ME 595 · University of Washington · Spring 2026

Abstract

Safety-critical autonomous systems require controllers that can be formally verified, audited, and deployed on resource-constrained hardware — properties that large neural networks cannot satisfy. SINDy-RL [1] addresses this by co-training a polynomial surrogate and a neural policy in a Dyna loop, then distilling the result into a closed-form polynomial. We stress-test this pipeline on the inverted double pendulum (IDP) — a two-link system with two coupled unstable modes — asking whether it delivers data-efficiency, sparsity, and interpretability simultaneously. After resolving two non-obvious engineering obstacles (polynomial degree ceiling and surrogate exploitation), the Dyna loop converges in six iterations using 39,616 real-environment steps — **10.1× fewer than a full-order PPO baseline** — with the best checkpoint achieving 90% task success. Behavioral cloning of this policy into a degree-3 polynomial ($R^2 = 0.989$) fully preserves 90% closed-loop success, but achieves only minimal sparsification: STLSQ drops just 5 of 165 possible terms, leaving the polynomial near-dense. This is not a threshold-tuning problem — a 500× sweep of the sparsity threshold confirms performance holds at 90–95% success as term count falls from 165 to 121, but the polynomial never approaches the compact, auditable form the method promises. The bottleneck is the polynomial feature basis itself, which is poorly conditioned: further zeroing risks eliminating structurally important contributions rather than numerical noise. A physics-informed variant exploiting IDP translational symmetry achieves comparable sparsity at reduced success (85%), confirming that the density is a property of this representation class on this system rather than an artifact of data volume or threshold choice.

1 Introduction

1.1 Motivation: The Inspectability Requirement

Deployed autonomous systems face a constraint that capable controllers must also be *inspectable*. Formal standards such as DO-178C (avionics software) and IEC 62443 (industrial control) require analyzable, auditable control laws [2], [3]; surgical robotics regulators may require that a control law be certifiable before permitting autonomous maneuvers near tissue; embedded actuators on spacecraft or small aerial vehicles have no floating-point stack capable of running a neural network at control rates. A ten-thousand-parameter neural network — however capable — offers no handle for stability proofs or formal verification, and its memory footprint alone disqualifies it from microcontroller deployment. A closed-form polynomial is the opposite: each term has a physical interpretation, stability arguments can be constructed analytically, and a 165-term policy fits in kilobytes and evaluates as a single dot product.

This gap between what learning produces and what safety-critical systems can accept motivates a growing body of work on inherently interpretable controllers [3] — models whose structure is transparent by construction, not approximated after the fact.

1.2 SINDy and the Data Problem for Unstable Systems

Sparse Identification of Nonlinear Dynamics (SINDy [4]) offers a principled path to interpretable governing equations. Given a library of candidate functions over the state and input, sparse regression identifies which terms actually drive the dynamics and discards the rest. The resulting model is compact, physically grounded, and composed of a small number of terms that a practitioner can read and reason about. The difficulty is data. For an unstable equilibrium — such as an inverted pendulum — a random policy crashes in a handful of steps, and the near-upright transitions that SINDy needs are entirely unvisited. The system cannot provide the training data without a controller it does not yet have.

The resolution is to co-train the dynamics model and the controller iteratively. Sutton’s Dyna architecture [5] alternates between training a policy in a learned surrogate model and collecting new data from the real environment, so each iteration improves both components. Zolman et al. [1] apply this principle to SINDy by using an ensemble SINDy model as the Dyna surrogate, yielding SINDy-RL — a framework that bootstraps the data problem while retaining interpretability as a downstream option.

1.3 Goals and Contributions

We stress-test SINDy-RL against its three core promises — data-efficient, sparse, and interpretable — on the inverted double pendulum (IDP), a demanding two-link benchmark with two coupled unstable modes and a narrow region of attraction. Our contributions are:

1. **Data efficiency confirmed.** The Dyna loop converges in six iterations using 39,616 real-environment steps — **10.1× fewer than a full-order PPO baseline** — with the best checkpoint achieving 90% task success.
2. **Near-sparsity via behavioral cloning.** STLSQ distillation of the converged neural policy yields a **160/165-term degree-3 polynomial** ($\lambda = 0.10$) that fully preserves 90% closed-loop success; a threshold ablation confirms robustness across a $500\times$ range of λ values (90–95% success, 142–165 terms).
3. **Root cause of residual density.** The polynomial feature matrix is ill-conditioned ($\kappa \approx 2.4 \times 10^4$), preventing STLSQ from zeroing additional terms without risking structural contributions — a representation issue, not a data-volume problem.
4. Diagnosis and resolution of two non-obvious engineering obstacles (polynomial degree ceiling and surrogate exploitation) with practical safeguards applicable to other unstable benchmark systems.

1.4 Ethics and Safety Considerations

Benefits. Interpretable controllers have direct safety benefits: closed-form polynomial expressions can be analyzed for failure modes before deployment, admit Lyapunov-style stability certificates in certain regimes, and allow human engineers to audit the control law. SINDy-RL’s data efficiency reduces wear and risk on real hardware during training — fewer physical interactions means fewer dangerous exploration failures. Compact controllers enable deployment on hardware that cannot run large neural networks, potentially bringing safe autonomous control to resource-constrained applications in healthcare, agriculture, and infrastructure.

Risks. Surrogate exploitation — where a policy finds action sequences that the polynomial model predicts as highly rewarding but that do not correspond to real physics — is a failure mode that can produce overconfident deployment decisions. A practitioner who observes high surrogate reward without validating on the real environment may incorrectly conclude the system is ready for deployment. Incomplete sparsity (our 165-term result is compact but not truly auditable in the way a five-term equation would be) creates a risk that the “interpretability” framing overstates what an engineer can actually verify. Finally, distribution shift between the training distribution and deployment conditions can cause polynomial controllers to fail catastrophically in states not covered by the training data — a risk that perturbation augmentation mitigates but does not eliminate.

Any deployment of a SINDy-RL-derived controller in a safety-critical application should include real-hardware validation across the full intended operating envelope, formal analysis of the polynomial’s behavior at boundary conditions, and a fallback mechanism (e.g., a conservative linear controller) for states outside the validated region.

2 Technical Background

2.1 The Testbed: Inverted Double Pendulum

The `InvertedDoublePendulum-v5` environment (MuJoCo 3.8.1 / Gymnasium 1.2.3) consists of two rigid links of equal length $L_1 = L_2 = 0.6$ m mounted on a sliding cart. The physical state is $\mathbf{x} = [x, \theta_1, \theta_2, \dot{x}, \dot{\theta}_1, \dot{\theta}_2] \in \mathbb{R}^6$, where x is the cart’s horizontal position along the track, θ_1, θ_2 are joint angles measured from vertical, and dots denote time derivatives. The 9-dimensional observation replaces raw angles with sine/cosine encodings to avoid wrapping discontinuities. The single control input is a horizontal cart force $u \in [-1, 1]$.

Tip height $h = L_1 \cos \theta_1 + L_2 \cos(\theta_1 + \theta_2)$ reaches a maximum of 1.2 m when both poles are vertical. Gymnasium terminates an episode when $h \leq 1.0$ m, leaving only a 0.2 m near-upright band between success and failure. The per-




Figure 1. IDP geometry. State $\mathbf{x} = [x, \theta_1, \theta_2, \dot{x}, \dot{\theta}_1, \dot{\theta}_2]$. Tip height $h \in [0, 1.2]$ m; episode ends at $h \leq 1.0$ m.

step reward is:

$$r_k = 10 \cdot \mathbf{1} - (h_k - 2)^2 - 0.01 x_{\text{tip}}^2 - \varepsilon \|\dot{\theta}\|^2$$

where the alive bonus (first term, $\approx 10/\text{step}$) dominates when the system remains upright. Episodes are capped at 1,000 steps (50 s at $\Delta t = 0.05$ s).

2.2 SINDy-C: Sparse Dynamics Identification with Control

SINDy [4] identifies discrete-time dynamics by regressing the state increment against a polynomial library:

$$\mathbf{x}_{k+1} - \mathbf{x}_k = \underbrace{\Theta(\mathbf{x}_k, u_k)}_{\text{library}} \cdot \underbrace{\Xi}_{\text{sparse coefficients}}$$

For control-affine systems (SINDy-C [6]), the input u_k is included directly in the library.¹ The Sequentially Thresholded Least Squares (STLSQ) algorithm zeros coefficients below threshold λ , promoting sparsity in Ξ . A degree- d library over n variables contains $\binom{n+d}{d}$ terms; for the IDP’s 7-dimensional state-action vector, degree-2 gives 36 features and degree-3 gives 120, a distinction that proved critical to convergence (§4.2).

2.3 E-SINDy: Ensemble Uncertainty Quantification

A single SINDy model provides a point estimate with no uncertainty information. Fasel et al. [7] address this with Ensemble SINDy (E-SINDy): fit M independent SINDy models on 80% bootstrap subsamples of the data, then at inference time report the mean and standard deviation of predictions across the ensemble. For $M = 10$ models, at each surrogate step $\text{predict}(\mathbf{x}, \mathbf{u})$ returns $(\mu_\Delta, \sigma_\Delta)$, where μ_Δ is the ensemble-mean predicted state increment and σ_Δ is the per-component standard deviation across members. High σ_Δ signals extrapolation beyond the training distribution. Following Zolman et al. §3.5 [1], we convert this into an active penalty: surrogate reward is reduced by $\kappa \cdot \text{mean}(\sigma_\Delta)$ per step ($\kappa = 5.0$), steering PPO away from high-uncertainty states.

2.4 Dyna-Style Model-Based RL and Behavioral Cloning

The Dyna architecture [5] alternates cheap model-based rollouts inside a learned surrogate with real-environment data collection. In SINDy-RL [1], the surrogate is the E-SINDy ensemble and the planner is PPO [8]. Figure 2 shows the RL control loop; in SINDy-RL the environment is instantiated twice: as the E-SINDy surrogate for cheap policy training, and as real MuJoCo for data collection and evaluation.




Figure 2. The RL control loop. The agent outputs action $u_k = \pi_\phi(\mathbf{x}_k)$ from the neural policy π_ϕ during Dyna training; after distillation into a polynomial, this becomes $u_k \approx \Theta_{\text{obs}}(\mathbf{x}_k) \xi$ where ξ is the sparse coefficient vector. In SINDy-RL, the environment is instantiated as the E-SINDy polynomial surrogate for cheap policy training and as the full MuJoCo simulator for real data collection and evaluation.

A Schroeder multi-sine sweep [9] bootstraps the initial dataset $\mathcal{D}$. Each Dyna iteration refits E-SINDy on near-upright transitions, runs PPO for 100k surrogate steps (warm-started from the prior policy), then collects 4,000 real transitions. After convergence, the best checkpoint is distilled via behavioral cloning:

¹A system is control-affine if the control input appears linearly in the dynamics: $\dot{\mathbf{x}} = f(\mathbf{x}) + g(\mathbf{x})u$, where f and g may be arbitrarily nonlinear in the state. Most mechanical systems driven by forces or torques, including the IDP, satisfy this property.

$$\min_{\Xi} \|\Theta_{\text{obs}}(X) \Xi - U^*\|_2 \quad (\text{STLSQ}, \lambda = 0.10)$$

where X is a matrix of observations, U^* are the corresponding NN policy actions, and Θ_{obs} is the degree-3 polynomial library over the 8-dimensional sin/cos observation. Perturbation augmentation [10] (adding Gaussian noise to expert states and re-querying the trained neural network policy, referred to as the NN oracle) expands the 50k-transition dataset $5\times$ to mitigate distribution shift without additional simulator rollouts.




Figure 3. Sparse regression structure for E-SINDy dynamics identification (*left*) and policy distillation (*right*). Each of the $k = 1, \dots, M$ ensemble members solves $\Delta \mathbf{X}^k = \Theta_{\text{dyn}}^k \Xi_{\text{dyn}}^k$, where $\Theta_{\text{dyn}}^k \in \mathbb{R}^{N \times 120}$ is the degree-3 polynomial library over the 7-dimensional state-action input evaluated on a bootstrap subsample, and $\Xi_{\text{dyn}}^k \in \mathbb{R}^{120 \times 6}$ (orange) is the fitted coefficient matrix. Distillation fits a single $U = \Theta_{\pi} \Xi_{\pi}$, where $\Theta_{\pi} \in \mathbb{R}^{N \times 165}$ is the degree-3 library over the 8-dimensional sin/cos observation, and $\Xi_{\pi} \in \mathbb{R}^{165 \times 1}$ (purple) is the scalar action coefficient vector. The two libraries are distinct: different input spaces (raw state-action vs. sin/cos encoding) and different feature counts (120 vs. 165). Annotations below each Ξ give the nonzero coefficient count after STLSQ thresholding — 690 of 720 possible entries for the dynamics model and 160 of 165 for the distilled policy — indicating that both fits are near-dense.

3 Methods

The evaluation has three components. First, a full-order PPO agent trained with unlimited simulator access establishes the performance ceiling. Second, the SINDy-RL Dyna pipeline co-trains an E-SINDy surrogate and a PPO policy using a fraction of those real interactions, targeting comparable task success. Third, behavioral cloning — fitting a sparse polynomial to imitate the neural policy’s actions, a process we refer to as *distillation* — converts the result into a closed-form expression suitable for embedded deployment and formal analysis. All three are implemented on the IDP testbed (§2.1).

3.1 Baseline: Full-Order PPO

The performance ceiling is a standard PPO agent (Stable-Baselines3 2.8.0) trained with unlimited real-environment access: a two-hidden-layer [64,64] MLP with tanh activations (9,731 parameters), trained for 400,000 total steps across 15,103 episodes. This agent is evaluated only as a reference point; the Dyna loop produces its own neural policy, and it is that policy — not the baseline — that is subsequently distilled into a polynomial.

3.2 SINDy-RL Pipeline

All code is implemented in Python 3.12.7 using PySINDy 2.1.0 [11] (E-SINDy surrogate), Stable-Baselines3 2.8.0 [12] (PPO policy), Gymnasium 1.2.3 with MuJoCo 3.8.1 [13] (simulation environment), NumPy 2.4.6 (numerical operations), and scikit-learn 1.8.0 (polynomial feature generation). The full pipeline is in `notebooks/sindy-rl.ipynb`.

Bootstrap (Stage 1). 300 episodes of Schroeder multi-sine excitation collect 2,897 near-upright state-transition pairs from real MuJoCo.

E-SINDy fit (Stage 2). Transitions with tip height $h > 1.10$ m (poles within $\approx 24^\circ$ of vertical) are retained.² Ten degree-3 SINDy-C models (PySINDy SINDy with PolynomialLibrary(degree=3) and STLSQ(threshold=0.05)) are fit on 80% bootstrap subsamples. Their coefficient matrices are stacked into a FastEnsemblePredictor³ for efficient surrogate stepping.

Surrogate PPO (Stage 3). The predictor is wrapped in EnsembleSurrogateEnv⁴, a Gymnasium environment that replicates MuJoCo’s reward formula and $h \leq 1.0$ m termination condition, adds the uncertainty penalty $\kappa \cdot \text{mean}(\sigma_\Delta)$, and enforces physical state bounds ($|x| \leq 2.5$ m, $|\theta| \leq 0.9$ rad, $|\dot{\theta}| \leq 12$ rad/s). PPO trains for 100k steps; early-stopped if mean surrogate episode length stays below 5 steps after 50k steps.

Real data collection (Stage 4). The policy is deployed in real MuJoCo for 4,000 steps, appended to $\mathcal{D}$.

Repeat with warm-start (Stage 5). If exploitation is detected (surrogate reward $> 3 \times$ previous AND real episode length $< 50\%$ of best seen), the next iteration rolls back to the best real-env checkpoint.

Distillation. 50k expert transitions are collected from the best Dyna checkpoint in real MuJoCo, augmented $5 \times$ with per-dimension Gaussian noise (σ : [0.02, 0.02, 0.02, 0.05, 0.10, 0.10] for $[x, \theta_1, \theta_2, \dot{x}, \dot{\theta}_1, \dot{\theta}_2]$), and queried from the NN oracle. A degree-3 polynomial is then fit via STLSQ ($\lambda = 0.10$) on the 8-dimensional sin/cos observation.

3.3 Metrics

Data efficiency is measured by real-environment step count. Task performance is measured by success rate (episodes lasting at least 500 steps) and mean episode length. Surrogate quality is tracked by E-SINDy one-step RMSE on held-out near-upright transitions. Distillation quality is measured by OLS R^2 and STLSQ term count.

4 Results

4.1 Baseline

Full-order PPO achieves mean reward $9,324 \pm 2$, 100% success, and mean episode length 1,000/1,000 steps at a cost of 400,000 real simulator interactions and a 9,731-parameter opaque network. This is the performance ceiling.

4.2 Dyna Loop Convergence

The Dyna loop converged in six iterations using 39,616 real steps, **10.1 $\times$ fewer than the baseline**. Figure 3 shows episode length distributions for selected checkpoints. Cross-validated evaluation of all saved checkpoints identified **iteration 7** as the best policy: **90% success and mean episode length 904 steps**. The converged E-SINDy dynamics model is dense: 690 out of 720 possible coefficients (120 features $\times$ 6 state dimensions) are nonzero, indicating that the IDP requires the full expressiveness of a degree-3 polynomial.

Iteration	Cumul. real steps	SINDy RMSE	Surr. mean len	Real mean len	Success
Bootstrap	2,897	0.016	—	—	—
1	7,015	0.016	11.8	12	0%
2	11,186	0.020	17.1	17	0%

²The initial value was inherited from a reward-shaping constant `TIP_HEIGHT_TARGET = 2.0` (a dimensionless offset in MuJoCo’s reward formula, not a physical height), setting `SINDY_H_MIN = 1.6` m — above the physical maximum $L_1 + L_2 = 1.2$ m. Every iteration silently fell back to fitting on all collected data until corrected to `SINDY_H_MIN = 1.10` m, derived from segment geometry.

³The bottleneck was `PolynomialLibrary.transform()` from scikit-learn, which carries ~ 1 ms fixed overhead per call regardless of input size; calling it once per ensemble member cost ~ 10 ms/step and ~ 13 min per 75k-step PPO phase. The fix: at construction, extract the `powers_` exponent matrix from sklearn’s `PolynomialFeatures` once; at each step, compute features as `np.prod(xu ** powers_, axis=1)` in pure NumPy and apply all 10 pre-stacked (10, 6, 120) coefficient matrices via a single batched matmul (~ 0.93 ms/step, $11.5 \times$ speedup).

⁴Gymnasium wrapper around the E-SINDy ensemble: reward = MuJoCo formula (§2.1) minus $\kappa \cdot \text{mean}(\sigma_\Delta)$; termination at $h \leq 1.0$ m or outside physical limits with a -50 out-of-bounds penalty. Matching MuJoCo’s conditions exactly is necessary for policy transfer.

Iteration	Cumul. real steps	SINDy RMSE	Surr. mean len	Real mean len	Success
3	15,551	0.084	36.5	36	0%
4	19,979	0.095	42.8	43	0%
5	29,453	0.085	547	547	50%
6	39,616	0.080	616	616	60%

The loop declared convergence at iteration 6 (60% surrogate success); the best checkpoint was selected by cross-evaluating all saved policies in real MuJoCo and is at iteration 7 (one additional PPO phase beyond convergence), achieving 90% success. The RMSE rise at iterations 3–4 reflects a better policy exploring states further from vertical, not model degradation; the surrogate remained accurate in the near-upright band.

4.3 Engineering Obstacles

Two non-obvious obstacles prevented convergence on early attempts.

Degree-2 RMSE ceiling. Over 25 iterations with data growing from 5k to 90k transitions, RMSE oscillated at 0.10–0.16 and real episode length grew from only 6 to 22 steps. A degree-2 library (36 features) cannot express the inter-modal coupling terms that dominate IDP dynamics (e.g., $\cos \theta_1 \cdot \cos \theta_2 \cdot \dot{\theta}_1$), which are cubic. When RMSE fails to decrease with $10\times$ more data, the cause is model capacity, not data quantity. Fix: SINDY_DEGREE=3 (120 features), which dropped RMSE to 0.013 within two iterations.

Surrogate exploitation. In a diagnostic run, surrogate reward jumped $9\times$ (497 to 4,525) while real episode length collapsed 87% (414 to 56 steps). The policy found action sequences the polynomial predicted as highly rewarding that had no correspondence to real physics. Uncertainty penalization alone is insufficient: all 10 ensemble members share the same polynomial basis, so in extrapolated regions they all make the same wrong prediction simultaneously and ensemble disagreement σ_Δ remains low. Rollback alone is also insufficient: it detects exploitation after the fact but does not prevent the exploited iteration from degrading the dataset. Fix: both mechanisms together — uncertainty penalty `reward -= 5.0 * mean(sigma_delta)` during surrogate PPO, plus rollback to best real-env checkpoint when surrogate reward exceeds $3\times$ the previous value and real episode length drops below 50% of best seen. The design principle generalizes: any ensemble surrogate built on a shared function basis cannot detect shared extrapolation errors through internal disagreement — real-environment feedback is the only reliable out-of-distribution signal.

4.4 Policy Distillation

Behavioral cloning from the best Dyna checkpoint (iteration 7, 90% success) produced a degree-3 polynomial with $R^2 = 0.9894$. Three additional obstacles required resolution: (1) the distillation teacher must be the best-checkpoint policy, not the final loop policy, which may have drifted during surrogate training (using the final policy gave 0% success); (2) degree-2 gave $R^2 \approx 0.905$ regardless of data volume, requiring degree-3; (3) perturbation augmentation was needed to close distribution shift between the NN’s training trajectories and deployment states.

STLSQ at $\lambda = 0.10$ dropped five terms, retaining **160/165 terms** with **90% closed-loop success** — identical to the teacher (Figure 5). The distilled policy is $59\times$ smaller than the baseline NN. Figure 6 shows a threshold ablation across $\lambda \in [0.001, 1.0]$: success remains at 90–95% while term count drops from 165 to 121, confirming that distillation performance is robust to the threshold choice. The dominant policy terms include recognizable physical structure: a constant bias; proportional angle feedback on $\cos \theta_1$ and $\cos \theta_2$; velocity damping on $\dot{\theta}_1$ and $\dot{\theta}_2$; and a large inter-pole coupling term ($\cos \theta_1 \times \cos \theta_2$) — the physically expected dominant nonlinearity. The remaining cubic cross-terms encode residual nonlinearity from the NN’s tanh activations.

The dynamics feature matrix Θ ($209,620 \times 120$) has condition number $\kappa = 2.37 \times 10^4$ (full rank, 120/120 singular values nonzero). Ill-conditioning worsens with output dimension: $\kappa = 3.18 \times 10^2$ for the x -position equation (40 active OLS features) rising to $\kappa = 2.30 \times 10^4$ for $\dot{\theta}_2$ (116 active features). In an ill-conditioned feature basis, small OLS coefficients along low-singular-value directions are numerically unreliable; STLSQ




Figure 4. Episode length distributions for the baseline PPO and selected SINDy-RL checkpoints, evaluated over 20 episodes each. Success rate (≥ 500 steps) is annotated above each box. Iteration 7 (starred) is the cross-validated best checkpoint. Iterations 10 and 20 illustrate the performance collapse following peak convergence — a consequence of continued surrogate exploitation that degrades the dataset.




Figure 5. Distilled policy coefficient magnitudes (log scale), coloured by polynomial degree. Of 165 possible terms, 160 are nonzero at $\lambda = 0.10$; five terms are dropped. Performance is fully preserved (90% success).

cannot safely threshold the remaining 160 terms without risking structural contributions. The 690/720 nonzero dynamics coefficients reflect the same phenomenon — not that all terms are mechanistically necessary, but that the feature basis does not admit a stable sparse decomposition. As on other unstable systems, $R^2 = 0.9894$ on the distillation set is a necessary but not sufficient metric: small in-distribution approximation errors can compound against the IDP’s unstable modes in closed loop.

Distillation preserves neural-policy performance while delivering a closed-form controller — the trade-off is justified by downstream requirements (embedded hardware, formal verification, or regulatory auditability).

Approach	Real-env steps	Mean ep len	Success	Params
Baseline PPO	400,000	1,000	100%	9,731
SINDy-RL NN (Dyna)	39,616	904	90%	9,731
SINDy-RL Polynomial	89,616 [†]	904	90%	160 terms

[†]\$39,616 Dyna steps + 50,000 rollout steps for distillation data; $5\times$ perturbation augmentation reuses the NN oracle without additional MuJoCo interactions.




Figure 6. STLSQ threshold ablation. Left axis: mean episode length (blue); right axis: non-zero term count (orange). Dashed vertical line marks the report threshold ($\lambda = 0.10$). Performance is stable at 90–95% success across the full ablation range, even as the polynomial compresses from 165 to 121 terms.

4.5 Physics-Informed Feature Reduction via Translational Symmetry

The IDP equations of motion are invariant under cart translations $x \rightarrow x + c$: the Lagrangian contains $\dot{x}$ but not x , so the state increment $\Delta x = \dot{x} \Delta t$ is given exactly by kinematics and need not be learned. Removing x from the SINDy library reduces the state-action input from 7 to 6 dimensions and cuts the degree-3 feature count from $\binom{10}{3} = 120$ to $\binom{9}{3} = 84$ — a 30% reduction — while hard-coding $\Delta x = \dot{x} \cdot \Delta t$ in the surrogate step. The distillation library is unchanged (it operates on the 8-dimensional sin/cos observation, which retains x via the position-dependent reward).

With the reduced library the Dyna loop converges in seven iterations using approximately 31,000 real-environment steps. The converged neural policy achieves **85% task success and mean episode length 856 steps**. Behavioral cloning yields $R^2 = 0.9987$ and **162/165 terms** at $\lambda = 0.05$, comparable to the standard formulation’s 160/165 at $\lambda = 0.10$, but at 5 percentage points lower neural-policy success (85% vs. 90%). The no- x variant does not outperform the standard formulation: both achieve comparable distillation sparsity while the standard formulation achieves higher task success. This suggests the degree-3 polynomial basis already captures the translational symmetry implicitly through the relationship between x , $\dot{x}$, and the dynamics, making the explicit coordinate removal redundant. We also examined SE(3) forward-kinematics coordinates (replacing raw angles with absolute-link sin/cos), which produced catastrophic ill-conditioning ($\kappa = 10^{17}$ – 10^{19}) due to unit-circle constraints ($\sin^2 \theta + \cos^2 \theta = 1$) creating near-exact linear dependence in the degree-2+ polynomial features — confirming that the choice of coordinate representation has a larger effect on numerical stability than on task performance, and that degree-3 in raw angles is difficult to improve upon at this data scale.

4.6 Code Repository

All code and results: https://github.com/falconeaj1/ME_595. Key notebooks: `full-order-simulation.ipynb` (baseline PPO), `sindy-rl.ipynb` (SINDy-RL pipeline), and `sindy-rl-no-x.ipynb` (translational-symmetry variant). Professor Michelle Hickner added as collaborator (GitHub: mhickner).

5 Summary

We stress-tested SINDy-RL against its three core promises on the inverted double pendulum. **Data efficiency is confirmed:** the Dyna loop trains a 90%-success neural policy using 39,616 real-environment steps — $10.1\times$ fewer than the full-order PPO baseline. **Near-sparsity is achieved through behavioral cloning:** STLSQ distillation at $\lambda = 0.10$ retains 160/165 terms while fully preserving 90% closed-loop success, and a threshold ablation confirms robustness across a $500\times$ range of threshold values (90–95% success, 121–165 terms). **Full sparsity remains elusive:** the 160 surviving terms are explained by ill-conditioning of the degree-3 feature matrix ($\kappa \approx 2.4 \times 10^4$) — a representation issue that neither more data nor aggressive thresholding can resolve without risking structural contributions. A 160-term polynomial is compact and evaluates as a single dot product, but is not auditable in the way a five-term equation would be.

Reaching the data-efficiency result required resolving two non-obvious engineering obstacles absent from Zolman et al.’s algorithm description: the degree-2 RMSE ceiling showed that the IDP requires a cubic polynomial library (120 features vs. 36); and surrogate exploitation showed that uncertainty penalization and rollback are complementary safeguards — neither is sufficient alone.

A physics-informed variant dropping cart position from the SINDy library (exploiting translational symmetry) achieves comparable distillation sparsity (162/165 at $\lambda = 0.05$) but at 5 percentage points lower neural-policy success (85% vs. 90%). The standard degree-3 library already captures the translational structure implicitly, suggesting that additional coordinate engineering is unlikely to be the path to true sparsity on this system.

Several open questions remain. First, it is unclear whether SINDy-RL can eventually match the baseline’s 100% success with fewer than 400,000 real steps if the Dyna loop is continued beyond the performance peak before exploitation sets in. Second, physics-informed libraries derived from the Euler-Lagrange equations — using gravity terms, mass-matrix coupling, and Coriolis interactions as basis functions rather than generic polynomial monomials — may achieve true sparsity by encoding the system’s structure directly; preliminary library comparisons show $35\times$ better conditioning ($\kappa \approx 3 \times 10^4$) with only 32 features versus the degree-3 polynomial’s 120, and a full Dyna comparison is ongoing. Third, the distillation process requires real-environment data collection despite the theoretical justification for querying the policy at any state [1]: “because there is no temporal dependence on π_ϕ , we can assemble our data and label pairs by evaluating π_ϕ for any $\mathbf{x}$ ” — it is unclear why surrogate trajectories alone are insufficient for distillation, and this bears further investigation.

CRedit Statement (<https://credit.niso.org>)

Role	Patrick Smith	Andrew Falcone
Conceptualization	Yes	Yes
Data curation	Yes	Yes
Formal analysis	Lead	Supporting
Investigation	Yes	Yes
Methodology	Lead	Supporting
Software	Lead	Supporting
Validation	Yes	Yes
Visualization	Yes	Yes
Writing – original draft	Yes	Yes
Writing – review & editing	Yes	Yes

AI tool disclosure: Claude (Anthropic) assisted with code drafting, debugging, writing iteration, and figure generation. All analysis, results, and conclusions were reviewed and executed by the authors, who take full responsibility for the submitted work.

References

- [1] N. Zolman, C. Lagemann, U. Fasel, J. N. Kutz, and S. L. Brunton, “SINDy-RL: Interpretable and efficient model-based reinforcement learning,” *Nature Communications*, 2025.
- [2] A. B. Arrieta *et al.*, “Explainable artificial intelligence (XAI): Concepts, taxonomies, opportunities and challenges toward responsible AI,” *Information Fusion*, vol. 58, pp. 82–115, 2020.
- [3] C. Rudin, “Stop explaining black box machine learning models for high stakes decisions and use interpretable models instead,” *Nature Machine Intelligence*, vol. 1, no. 5, pp. 206–215, 2019.
- [4] S. L. Brunton, J. L. Proctor, and J. N. Kutz, “Discovering governing equations from data by sparse identification of nonlinear dynamical systems,” *Proceedings of the National Academy of Sciences*, vol. 113, no. 15, pp. 3932–3937, 2016.
- [5] R. S. Sutton, “Integrated architectures for learning, planning, and reacting based on approximating dynamic programming,” in *International conference on machine learning (ICML)*, 1990, pp. 216–224.
- [6] E. Kaiser, J. N. Kutz, and S. L. Brunton, “Sparse identification of nonlinear dynamics for model predictive control in the low-data limit,” *Proceedings of the Royal Society A*, vol. 474, no. 2219, p. 20180335, 2018.
- [7] U. Fasel, J. N. Kutz, B. W. Brunton, and S. L. Brunton, “Ensemble-SINDy: Robust sparse model discovery in the low-data, high-noise limit,” *Proceedings of the Royal Society A*, vol. 478, no. 2260, 2022.
- [8] J. Schulman, F. Wolski, P. Dhariwal, A. Radford, and O. Klimov, “Proximal policy optimization algorithms.” 2017.
- [9] M. R. Schroeder, “Synthesis of low-peak-factor signals and binary sequences with low autocorrelation,” *IEEE Transactions on Information Theory*, vol. 16, no. 1, pp. 85–89, 1970.
- [10] S. Ross, G. Gordon, and D. Bagnell, “A reduction of imitation learning and structured prediction to no-regret online learning,” in *AISTATS*, in PMLR, vol. 15. 2011, pp. 627–635.
- [11] B. M. de Silva, K. Champion, M. Quade, J.-C. Loiseau, J. N. Kutz, and S. L. Brunton, “PySINDy: A Python package for the sparse identification of nonlinear dynamical systems from data,” *Journal of Open Source Software*, vol. 5, no. 49, p. 2104, 2020.
- [12] A. Raffin, A. Hill, A. Gleave, A. Kanervisto, M. Ernestus, and N. Dormann, “Stable-Baselines3: Reliable reinforcement learning implementations,” *Journal of Machine Learning Research*, vol. 22, no. 268, pp. 1–8, 2021.
- [13] E. Todorov, T. Erez, and Y. Tassa, “MuJoCo: A physics engine for model-based control,” in *IEEE/RSJ international conference on intelligent robots and systems (IROS)*, 2012, pp. 5026–5033.